CSC 522 : Automated Learning and Data Analysis

Homework 2

Roopak Venkatakrishnan - rvenkat7@ncsu.edu

January 27, 2013

1 Question 1

Write code in R or Matlab to perform each of the following tasks:

1. Generate a 3×3 matrix with input containing the sequence 1, 2, ... 9.
   Ans: $x \leftarrow matrix(c(1:9), 3, 3)$

2. a) Access elements from the 2nd and 3rd columns only.
       Ans:
       2nd column alone : x[,2] $\implies$ [1] 4 5 6
       3rd column alone : x[,3] $\implies$ [1] 7 8 9
       both columns : x[,2:3]
       $\implies$

	[,1]	[,2]
[1,]	4	7
[2,]	5	8
[3,]	6	9

    b) Access elements of the 2nd and 3rd rows only
       Ans:
       2nd row alone : x[2,] $\implies$ [1] 2 5 8
       3rd row alone : x[3,] $\implies$ [1] 3 6 9
       both rows : x[2:3,]
       $\implies$

	[,1]	[,2]	[,3]
[1,]	2	5	8
[2,]	3	6	9

    c) Access rows 1 and 3 only? (see rbind() function in R and vertcat() in matlab)
       Ans:
       $x2 \leftarrow rbind(x[2,], x[3,])$
       $\implies$

	[,1]	[,2]	[,3]
[1,]	1	4	7
[2,]	3	6	9

    d) Calculate sum of the 2nd row, the diagonal and the 3rd column in the matrix.
       Ans:
       sum(x[2,] + x[,3] + diag(x))
       $\implies$ [1] 54

e) Identify row and column dimensions of the matrix.
   Ans:
   dim(x)
   $\implies$ [1] 3 3

f) Transpose of a matrix.
   Ans:
   t(x)

	[, 1]	[, 2]	[, 3]
[1,]	1	2	3
[2,]	4	5	6
[3,]	7	8	9

g) Scalar multiplication of output matrix with itself.
   Ans:
   x * x

	[, 1]	[, 2]	[, 3]
[1,]	1	16	49
[2,]	4	25	64
[3,]	9	36	81

h) Matrix multiplication of output matrix with itself.
   Ans:
   x %*% x

	[, 1]	[, 2]	[, 3]
[1,]	30	66	102
[2,]	36	81	126
[3,]	42	96	150

i) Cross product of the output matrix from 1.
   Ans:
   crossprod(x)

	[, 1]	[, 2]	[, 3]
[1,]	14	32	50
[2,]	32	77	122
[3,]	50	122	194

j) Check if a matrix is a square matrix.
   Ans:

```
function checksqmatrix(mat)
{
  if(dim(mat)[1]==dim(mat)[2])
  {
    print("It is a square matrix!")
  }
  else
  {
    print("It is NOT a square matrix")
  }
}

> checksqmatrix(x)
[1] "It is a square matrix!"

> checksqmatrix(matrix(c(1:10),2,5))
[1] "It is NOT a square matrix"
```

k) Inverse of a matrix
   Ans:
   solve(x)
   Since this matrix has determinant 0 the inverse is not defined.
   Error in solve.default(x) :
   Lapack routine dgesv: system is exactly singular: U[3,3] = 0

l) Identity of a matrix.
   Ans:

   ```
   > x %*% solve(x)
   ```

m) Sum of all elements in the matrix (use a for/while loop)

   ```
   matrixsum <- function (mat) {
     i<-1
     sum<-0
     while(i<=dim(mat)[1]*dim(mat)[2])
     {
       sum<-sum+mat[i]
       i<-i+1
     }
     print (paste(sum, "is the sum of elements"))
   }

   > matrixsum(x)
   [1] "45 is the sum of elements"
   ```

2 Question 2

For this exercise, use the values.txt file provided. The file contains a list of 150 data instances. There are 2 columns representing the x and y coordinates. Complete the following tasks:

1. Load the file **Ans:**
   $\implies$
   ```
   > vals = read.table("D:\\Courses\\datamining - CSC522\\homework\\hw2\\values.txt")
   ```

2. Make a 2-D plot and label the axes
   plot(vals,main="A 2D plot of values.txt",xlab="Values of V1",ylab="Values of V2")
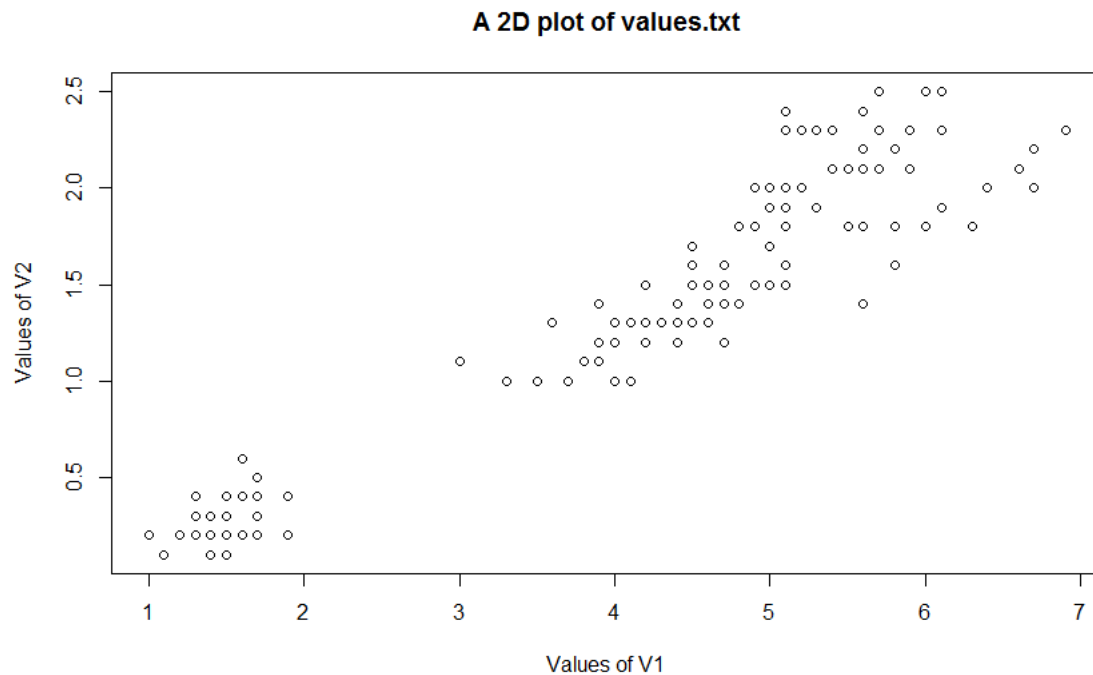
A 2D plot of values.txt

Figure 1: The 2D plot obtained from values.txt

3. Find the correlation between the dimensions.

```
> round(cor(vals3),3)
      V1    V2
V1 1.000 0.963
V2 0.963 1.000
```

4. Now consider a point (5; 1:5).

   a) Compute the distance of this point from each of the 150 instances using Euclidean distance, Mahalanobis distance, City block metric, Minkowski metric, Chebychev distance and Cosine distance.

   Ans:

   ### Euclidean Distance

```
library(fields)
rdist(vals,matrix(c(5,1.5),1,2))
              [,1]
  [1,] 3.8275318418
  [2,] 3.8275318418
  [3,] 3.9217343102
  [4,] 3.7336309405
  [5,] 3.8275318418
  [6,] 3.4785054262
  [7,] 3.7947331922
  [8,] 3.7336309405
  [9,] 3.8275318418
 [10,] 3.7696153650
 [11,] 3.7336309405
 [12,] 3.6400549446
```

```
[13,] 3.8626415832
[14,] 4.1436698710
[15,] 4.0162171256
[16,] 3.6687872656
[17,] 3.8600518131
[18,] 3.7947331922
[19,] 3.5114099732
[20,] 3.7000000000
[21,] 3.5468295702
[22,] 3.6687872656
[23,] 4.2059481690
[24,] 3.4481879299
[25,] 3.3615472628
[26,] 3.6400549446
[27,] 3.5735136770
[28,] 3.7336309405
[29,] 3.8275318418
[30,] 3.6400549446
[31,] 3.6400549446
[32,] 3.6687872656
[33,] 3.7696153650
[34,] 3.8275318418
[35,] 3.7336309405
[36,] 4.0162171256
[37,] 3.9217343102
[38,] 3.8626415832
[39,] 3.9217343102
[40,] 3.7336309405
[41,] 3.8897300678
[42,] 3.8897300678
[43,] 3.9217343102
[44,] 3.5171010790
[45,] 3.2893768407
[46,] 3.7947331922
[47,] 3.6400549446
[48,] 3.8275318418
[49,] 3.7336309405
[50,] 3.8275318418
[51,] 0.3162277660
[52,] 0.5000000000
[53,] 0.1000000000
[54,] 1.0198039027
[55,] 0.4000000000
[56,] 0.5385164807
[57,] 0.3162277660
[58,] 1.7720045147
[59,] 0.4472135955
[60,] 1.1045361017
[61,] 1.5811388301
[62,] 0.8000000000
[63,] 1.1180339887
[64,] 0.3162277660
[65,] 1.4142135624
[66,] 0.6082762530
[67,] 0.5000000000
[68,] 1.0295630141
[69,] 0.5000000000
[70,] 1.1704699911
```

```
[71,]   0.3605551275
[72,]   1.0198039027
[73,]   0.1000000000
[74,]   0.4242640687
[75,]   0.7280109889
[76,]   0.6082762530
[77,]   0.2236067977
[78,]   0.2000000000
[79,]   0.5000000000
[80,]   1.5811388301
[81,]   1.2649110641
[82,]   1.3928388277
[83,]   1.1401754251
[84,]   0.1414213562
[85,]   0.5000000000
[86,]   0.5099019514
[87,]   0.3000000000
[88,]   0.6324555320
[89,]   0.9219544457
[90,]   1.0198039027
[91,]   0.6708203932
[92,]   0.4123105626
[93,]   1.0440306509
[94,]   1.7720045147
[95,]   0.8246211251
[96,]   0.8544003745
[97,]   0.8246211251
[98,]   0.7280109889
[99,]   2.0396078054
[100,]  0.9219544457
[101,]  1.4142135624
[102,]  0.4123105626
[103,]  1.0816653826
[104,]  0.6708203932
[105,]  1.0630145813
[106,]  1.7088007491
[107,]  0.5385164807
[108,]  1.3341664064
[109,]  0.8544003745
[110,]  1.4866068747
[111,]  0.5099019514
[112,]  0.5000000000
[113,]  0.7810249676
[114,]  0.5000000000
[115,]  0.9055385138
[116,]  0.8544003745
[117,]  0.5830951895
[118,]  1.8384776311
[119,]  2.0615528128
[120,]  0.0000000001
[121,]  1.0630145813
[122,]  0.5099019514
[123,]  1.7720045147
[124,]  0.3162277660
[125,]  0.9219544457
[126,]  1.0440306509
[127,]  0.3605551275
[128,]  0.3162277660
```

```
[129,] 0.8485281374
[130,] 0.8062257748
[131,] 1.1704699911
[132,] 1.4866068747
[133,] 0.9219544457
[134,] 0.1000000000
[135,] 0.6082762530
[136,] 1.3601470509
[137,] 1.0816653826
[138,] 0.5830951895
[139,] 0.3605551275
[140,] 0.7211102551
[141,] 1.0816653826
[142,] 0.8062257748
[143,] 0.4123105626
[144,] 1.2041594579
[145,] 1.2206555616
[146,] 0.8246211251
[147,] 0.4000000000
[148,] 0.5385164807
[149,] 0.8944271910
[150,] 0.3162277660
```

Manhatten Distance

```
> abs(vals[["V2"]]-1.5) + abs(vals[["V1"]]-5)
  [1] 4.9 4.9 5.0 4.8 4.9 4.4 4.8 4.8 4.9 4.9 4.8 4.7
 [13] 5.0 5.3 5.1 4.6 4.8 4.8 4.5 4.7 4.6 4.6 5.3 4.3
 [25] 4.4 4.7 4.5 4.8 4.9 4.7 4.7 4.6 4.9 4.9 4.8 5.1
 [37] 5.0 5.0 5.0 4.8 4.9 4.9 5.0 4.3 4.2 4.8 4.7 4.9
 [49] 4.8 4.9 0.4 0.5 0.1 1.2 0.4 0.7 0.4 2.2 0.6 1.2
 [61] 2.0 0.8 1.5 0.4 1.6 0.7 0.5 1.4 0.5 1.5 0.5 1.2
 [73] 0.1 0.6 0.9 0.7 0.3 0.2 0.5 2.0 1.6 1.8 1.4 0.2
 [85] 0.5 0.6 0.3 0.8 1.1 1.2 0.9 0.5 1.3 2.2 1.0 1.1
 [97] 1.0 0.9 2.4 1.1 2.0 0.5 1.5 0.9 1.5 2.2 0.7 1.6
[109] 1.1 2.1 0.6 0.7 1.1 0.5 1.0 1.1 0.8 2.4 2.7 0.0
[121] 1.5 0.6 2.2 0.4 1.3 1.3 0.5 0.4 1.2 0.9 1.5 1.9
[133] 1.3 0.1 0.7 1.9 1.5 0.8 0.5 1.0 1.5 0.9 0.5 1.7
[145] 1.7 1.0 0.4 0.7 1.2 0.4
```

Cosine Distance

```
> cosnum <- vals[["V1"]]*5 + vals[["V2"]]*1.5
> cosden<- sqrt(27.25) * sqrt((vals[["V1"]]*vals[["V1"]])
+                             + (vals[["V2"]]*vals[["V2"]]))
> cosdist <- cosnum/cosden
> cosdist
  [1] 0.9888368 0.9888368 0.9903817 0.9874011 0.9888368 0.9981785 0.9967726
  [8] 0.9874011 0.9888368 0.9748189 0.9874011 0.9860710 0.9758649 0.9799079
 [15] 0.9920337 0.9995240 0.9999752 0.9967726 0.9931884 0.9955795 0.9848398
 [22] 0.9995240 0.9955795 0.9999854 0.9826444 0.9860710 0.9989201 0.9874011
 [29] 0.9888368 0.9860710 0.9860710 0.9995240 0.9748189 0.9888368 0.9874011
 [36] 0.9920337 0.9903817 0.9758649 0.9903817 0.9874011 0.9979104 0.9979104
 [43] 0.9903817 0.9977353 0.9964774 0.9967726 0.9860710 0.9888368 0.9874011
 [50] 0.9888368 0.9999981 0.9995412 0.9999843 0.9997407 0.9997178 0.9999477
```

```
 [57] 0.9993280 0.9999961 0.9998715 0.9985857 0.9999134 0.9986707 0.9989201
 [64] 0.9999981 0.9984834 0.9998623 0.9995412 0.9986366 0.9995412 0.9998631
 [71] 0.9977353 0.9997407 0.9999843 0.9991399 0.9999977 0.9998623 0.9999706
 [78] 0.9993419 0.9995412 0.9999134 0.9999531 0.9996221 0.9999752 0.9999213
 [85] 0.9995412 0.9987423 0.9998473 0.9999913 0.9998785 0.9997407 0.9996824
 [92] 0.9999921 1.0000000 0.9999961 0.9999620 0.9999134 0.9999620 0.9999977
 [99] 0.9982013 0.9998785 0.9946658 0.9978776 0.9987255 0.9998091 0.9974744
[106] 0.9998623 0.9975687 0.9999134 0.9999552 0.9952506 0.9966177 0.9986084
[113] 0.9973164 0.9960377 0.9890110 0.9930484 0.9996918 0.9996670 0.9995412
[120] 1.0000000 0.9957645 0.9953891 0.9999991 0.9981651 0.9981074 1.0000000
[127] 0.9977353 0.9981651 0.9977353 0.9997516 0.9999449 0.9999347 0.9965677
[134] 0.9999854 0.9989201 0.9976129 0.9935731 0.9996918 0.9977353 0.9968467
[141] 0.9935731 0.9912727 0.9978776 0.9967815 0.9925826 0.9921999 0.9974314
[148] 0.9971348 0.9938239 0.9988561
```

Chebyshev Distance

```
> v4<-cbind(abs(vals$V1-5),abs(vals$V2-1.5))
> tchebyshev <- sapply(v4,max)
> tchebyshev
  [1] 3.6 3.6 3.7 3.5 3.6 3.3 3.6 3.5 3.6 3.5 3.5 3.4 3.6 3.9 3.8
 [16] 3.5 3.7 3.6 3.3 3.5 3.3 3.5 4.0 3.3 3.1 3.4 3.4 3.5 3.6 3.4
 [31] 3.4 3.5 3.5 3.6 3.5 3.8 3.7 3.6 3.7 3.5 3.7 3.7 3.7 3.4 3.1
 [46] 3.6 3.4 3.6 3.5 3.6 0.3 0.5 0.1 1.0 0.4 0.5 0.3 1.7 0.4 1.1
 [61] 1.5 0.8 1.0 0.3 1.4 0.6 0.5 0.9 0.5 1.1 0.2 1.0 0.1 0.3 0.7
 [76] 0.6 0.2 0.0 0.5 1.5 1.2 1.3 1.1 0.1 0.5 0.5 0.3 0.6 0.9 1.0
 [91] 0.6 0.4 1.0 1.7 0.8 0.8 0.8 0.7 2.0 0.9 1.0 0.1 0.9 0.6 0.8
[106] 1.6 0.5 1.3 0.8 1.1 0.1 0.3 0.5 0.0 0.1 0.3 0.5 1.7 1.9 0.0
[121] 0.7 0.1 1.7 0.1 0.7 1.0 0.2 0.1 0.6 0.8 1.1 1.4 0.6 0.1 0.6
[136] 1.1 0.6 0.5 0.2 0.4 0.6 0.1 0.1 0.9 0.7 0.2 0.0 0.2 0.4 0.1
[151] 1.3 1.3 1.3 1.3 1.3 1.1 1.2 1.3 1.3 1.4 1.3 1.3 1.4 1.4 1.3
[166] 1.1 1.1 1.2 1.2 1.2 1.3 1.1 1.3 1.0 1.3 1.3 1.1 1.3 1.3 1.3
[181] 1.3 1.1 1.4 1.3 1.3 1.3 1.3 1.4 1.3 1.3 1.2 1.2 1.3 0.9 1.1
[196] 1.2 1.3 1.3 1.3 1.3 0.1 0.0 0.0 0.2 0.0 0.2 0.1 0.5 0.2 0.1
[211] 0.5 0.0 0.5 0.1 0.2 0.1 0.0 0.5 0.0 0.4 0.3 0.2 0.0 0.3 0.2
[226] 0.1 0.1 0.2 0.0 0.5 0.4 0.5 0.3 0.1 0.0 0.1 0.0 0.2 0.2 0.2
[241] 0.3 0.1 0.3 0.5 0.2 0.3 0.2 0.2 0.4 0.2 1.0 0.4 0.6 0.3 0.7
[256] 0.6 0.2 0.3 0.3 1.0 0.5 0.4 0.6 0.5 0.9 0.8 0.3 0.7 0.8 0.0
[271] 0.8 0.5 0.5 0.3 0.6 0.3 0.3 0.3 0.6 0.1 0.4 0.5 0.7 0.0 0.1
[286] 0.8 0.9 0.3 0.3 0.6 0.9 0.8 0.4 0.8 1.0 0.8 0.4 0.5 0.8 0.3
```

Minkowski Distance(p=3)

```
> vals2 <- cbind(abs(vals$V1-5),abs(vals$V2-1.5))
> minkowski <- (vals2[,1]^3 + vals2[,2]^3)^(1/3)
> minkowski
  [1] 3.6556427 3.6556427 3.7527387 3.5587893 3.6556427
  [6] 3.3402479 3.6439068 3.5587893 3.6556427 3.5731281
 [11] 3.5587893 3.4622056 3.6692360 3.9592317 3.8500534
 [16] 3.5358492 3.7321283 3.6439068 3.3520667 3.5464025
 [21] 3.3659226 3.5358492 4.0452569 3.3303295 3.1744052
 [26] 3.4622056 3.4379542 3.5587893 3.6556427 3.4622056
 [31] 3.4622056 3.5358492 3.5731281 3.6556427 3.5587893
 [36] 3.8500534 3.7527387 3.6692360 3.7527387 3.5587893
 [41] 3.7416049 3.7416049 3.7527387 3.4208921 3.1454962
```

```
[46] 3.6439068 3.4622056 3.6556427 3.5587893 3.6556427
[51] 0.3036589 0.5000000 0.1000000 1.0026596 0.4000000
[56] 0.5104469 0.3036589 1.7142970 0.4160168 1.1002754
[61] 1.5182945 0.8000000 1.0400419 0.3036589 1.4013592
[66] 0.6009245 0.5000000 0.9487518 0.5000000 1.1173556
[71] 0.3271066 1.0026596 0.1000000 0.3779763 0.7054004
[76] 0.6009245 0.2080084 0.2000000 0.5000000 1.5182945
[81] 1.2146356 1.3242015 1.1073883 0.1259921 0.5000000
[86] 0.5013298 0.3000000 0.6073178 0.9032802 1.0026596
[91] 0.6240251 0.4020726 1.0089202 1.7142970 0.8041452
[96] 0.8138223 0.8041452 0.7054004 2.0053192 0.9032802
[101] 1.2599210 0.4020726 0.9813199 0.6240251 0.9491220
[106] 1.6276446 0.5104469 1.3053038 0.8138223 1.3259101
[111] 0.5013298 0.4497941 0.6986368 0.5000000 0.9004113
[116] 0.8138223 0.5336803 1.7386752 1.9461462 0.0000000
[121] 0.9491220 0.5013298 1.7142970 0.3036589 0.8237661
[126] 1.0089202 0.3271066 0.3036589 0.7559526 0.8005205
[131] 1.1173556 1.4209436 0.8237661 0.1000000 0.6009245
[136] 1.2260507 0.9813199 0.5336803 0.3271066 0.6542133
[141] 0.9813199 0.8005205 0.4020726 1.0746258 1.1032959
[146] 0.8041452 0.4000000 0.5104469 0.8320335 0.3036589
```

b) For each distance measure, identify the 10 points from the dataset that are the closest to the point(5,1.5).

- Create plots, one for each distance measure. Place an X for (5,1.5) and mark the 10 closest points. To mark them, you could place a circle or any other shape over the point.

- Verify if the set of points is the same across all the distance measures.

.

.

Minkowski Distance(p=3)

```
> mink2 <- cbind(vals,minkowski)
> mink3<-head(mink2[order(minkowski),],10)
> plot(mink3[["V1"]],mink3[["V2"]],xlab="V1",ylab="V2",
+ main="Graph for Minkowski Distance",xlim=c(4.5,5.5),ylim=c(1.35,1.8))
> points(5,1.5,type="b",pch=4,col="#660000",cex=3)
```
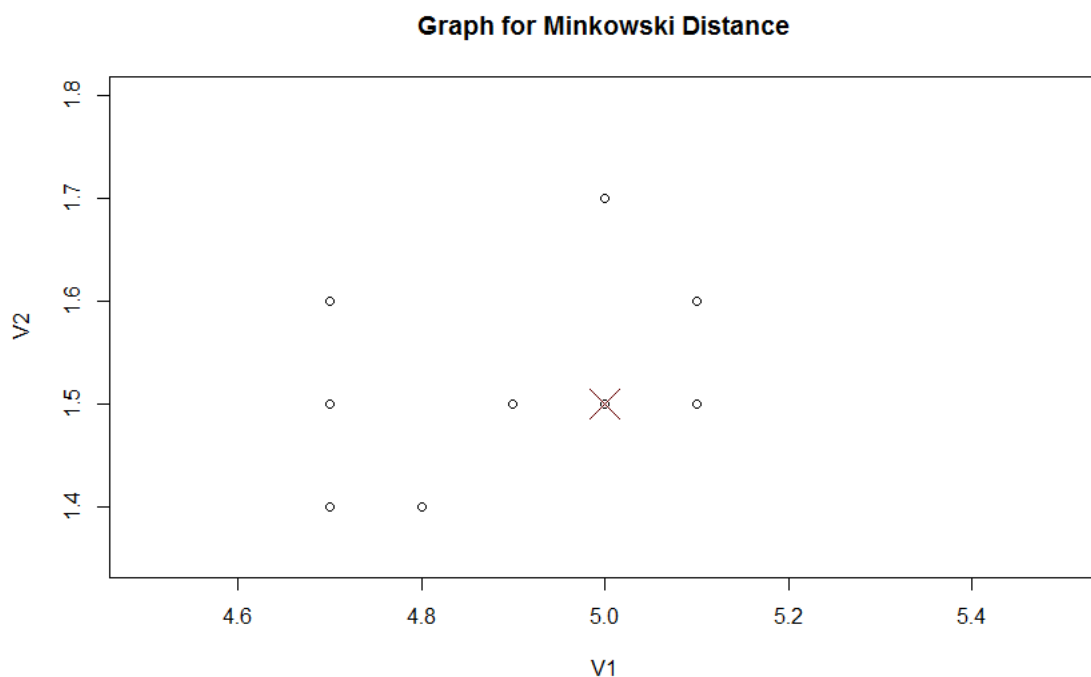

Figure 2: The plot for Minkowski Distance

Tchebyshev Distance

```
> tcheby1 <- cbind(vals,tchebyshev)
> tcheby2<-head(tcheby1[order(tchebyshev),],10)
> plot(tcheby2[["V1"]],tcheby2[["V2"]],xlab="V1",ylab="V2"
+ ,main="Graph for Tchebyshev Distance")
> points(5,1.5,type="b",pch=4,col="#660000",cex=3)
```
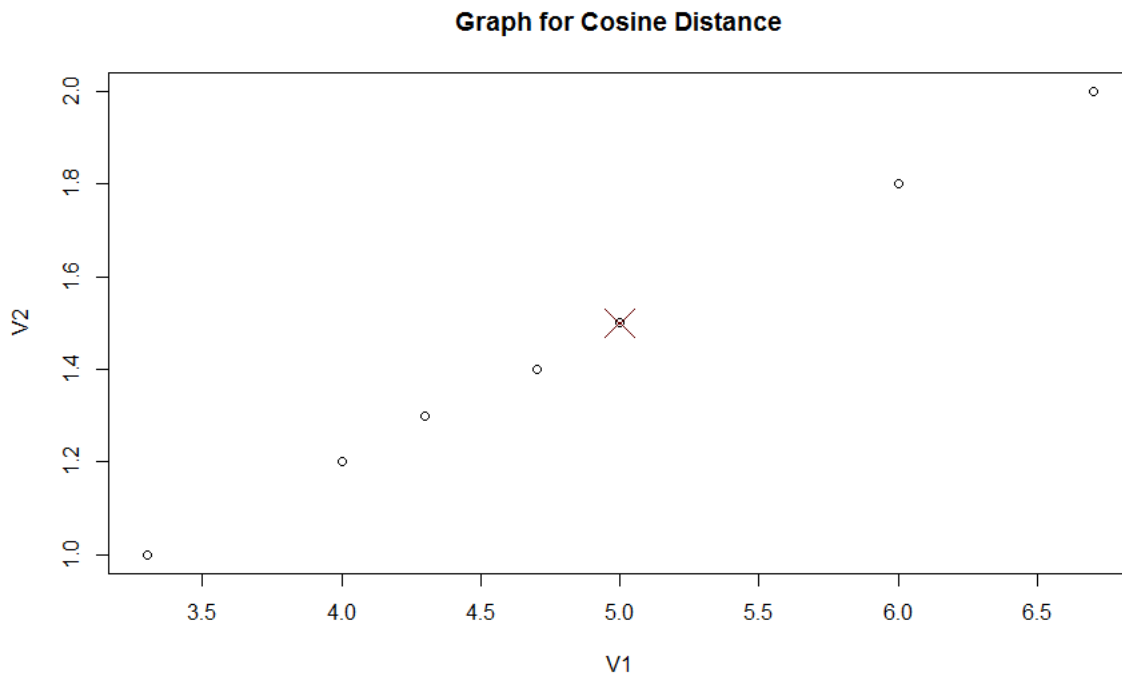
Graph for Tchebyshev Distance
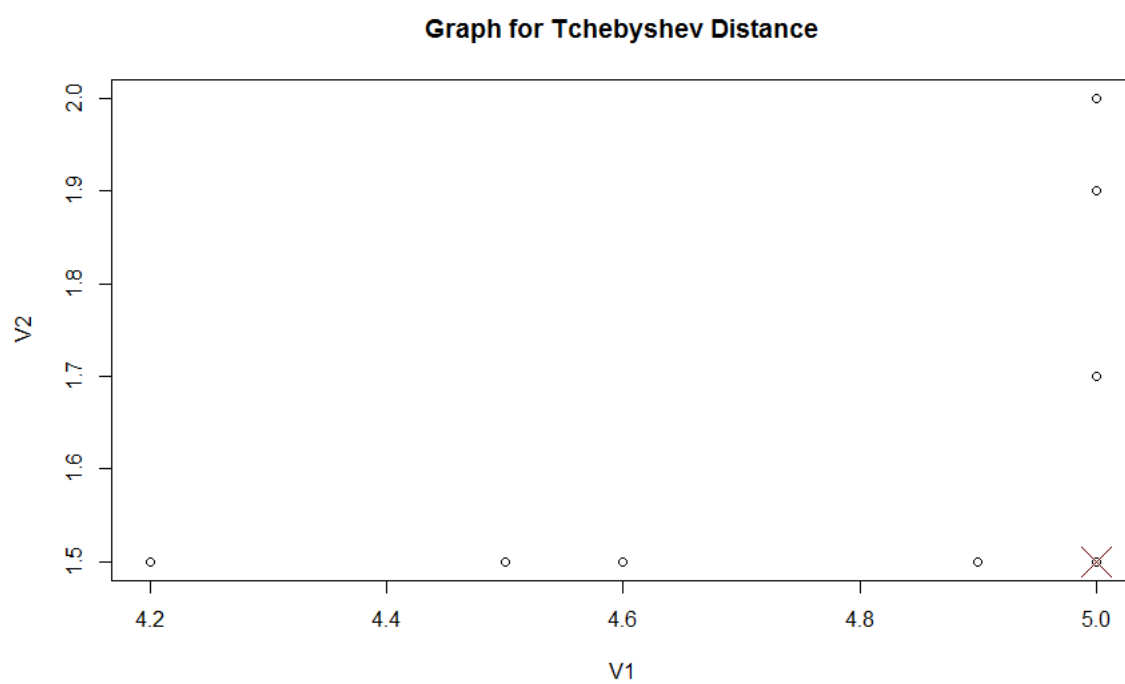
Figure 3: The plot for Tchebyshev Distance

Cosine Distance

Using Cosine Distance as opposde to cosine similarity.

```
> cos1 <- cbind(vals,cosdist)
> cos2<-head(cos1[order(cosdist),],10)
> plot(cos2[["V1"]],cos2[["V2"]],xlab="V1",ylab="V2",
+ main="Graph for Cosine Distance")
> points(5,1.5,type="b",pch=4,col="#660000",cex=3)
```

Graph for Cosine Distance

Figure 4: The plot for Cosine Distance

Manhattan Distance

```
> manh1 <- cbind(vals,manhattan)
> manh2<-head(manh1[order(manhattan),],10)
> plot(manh2[["V1"]],manh2[["V2"]],xlab="V1",ylab="V2"
+ ,main="Graph for Manhattan Distance")
> points(5,1.5,type="b",pch=4,col="#660000",cex=3)
```
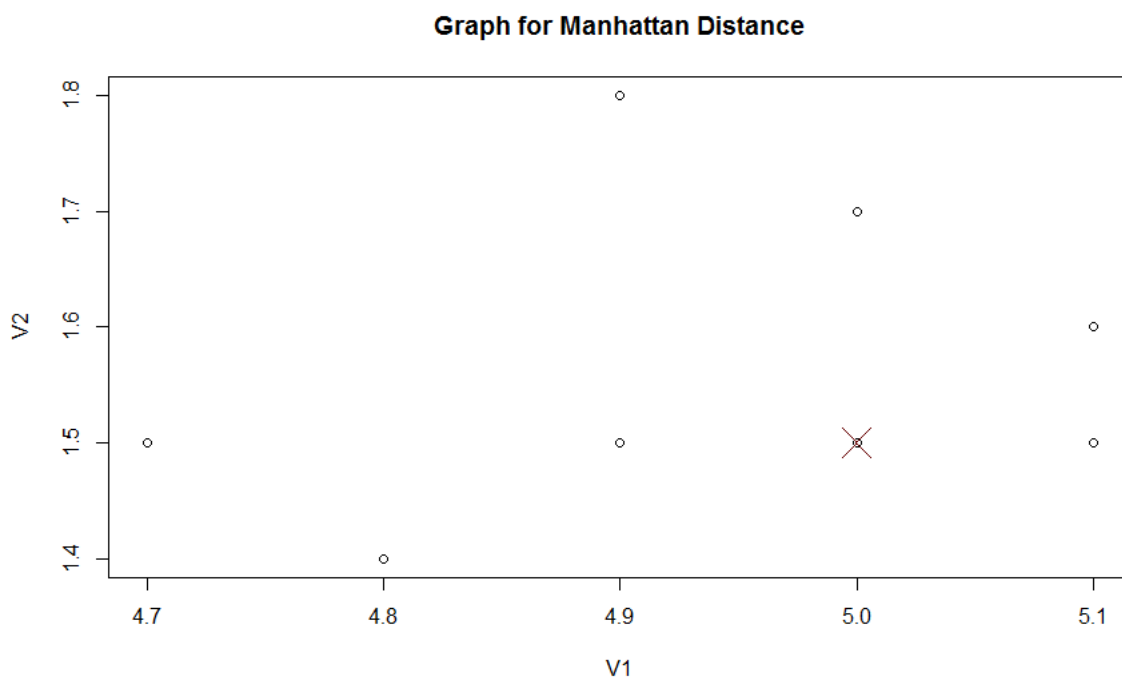

Figure 5: The plot for Manhattan Distance

Euclidean Distance

```
> euc1 <- cbind(vals,euclid)
> euc2<-head(euc1[order(euclid),],10)
> plot(euc2[["V1"]],euc2[["V2"]],xlab="V1",ylab="V2"
+ ,main="Graph for Euclidean Distance")
> points(5,1.5,type="b",pch=4,col="#660000",cex=3)
```

Graph for Euclidean Distance
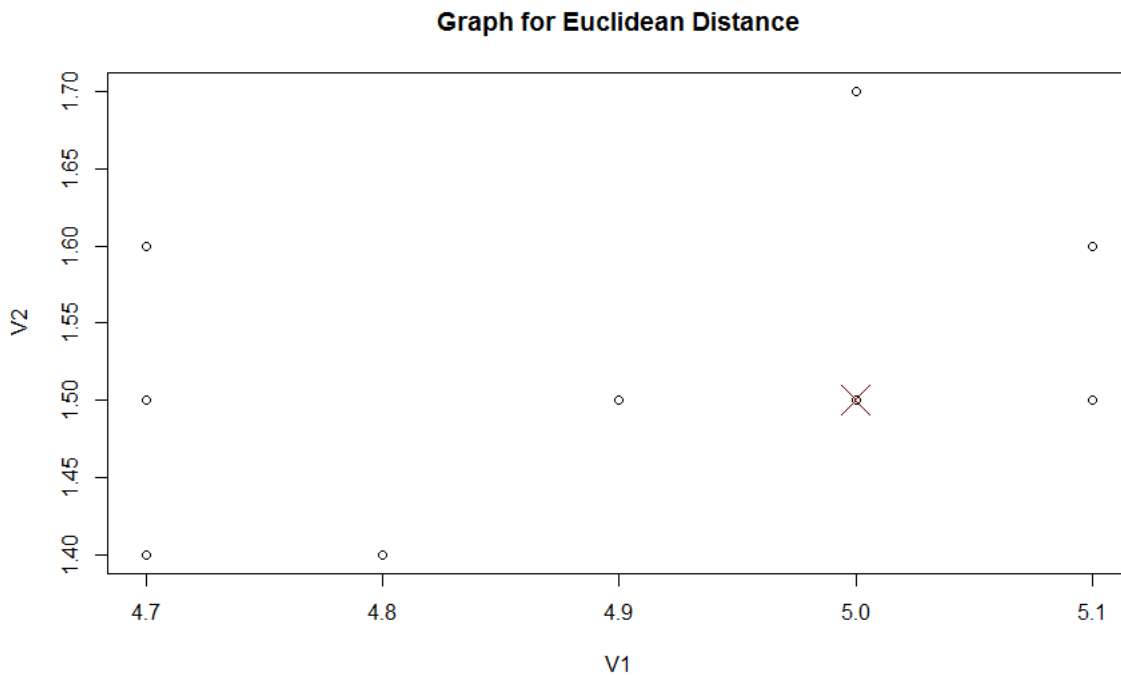
Figure 6: The plot for Manhattan Distance

3 Question 3

In this question you will be asked to summarize and explore data in file labelled hw1q3.csv, available on the course moodle site. This file contains 2 columns of 1000 data points each.

1. Load the data contained in the file using your platform of choice. Compute and report the mean, median, standard deviation, and the range (i.e. the minimum and maximum) for each variable x1, x2. As discussed in class, please be sure to use $s = \frac{1}{n-1} \sum_{t=1}^{n} (x_i - \bar{x})^2$ to compute the standard deviation. You dont have to do this by hand, but please verify any approach you use is based on this formulation. Report each estimate to 3 decimal places.

   - **Mean**

     ```
     > read.csv(file.choose(),header=T)
     > round(mean(vals$x1),digits=3)
     [1] 5.032
     > round(mean(vals$x2),digits=3)
     [1] 4.774
     ```

   - **Median**

```
> round(median(vals$x1),digits=3)
[1] 5.052
> round(median(vals$x2),digits=3)
[1] 2.567
```

- **Std. Deviation**

```
> round(sd(vals$x1),digits=3)
[1] 1.464
> round(sd(vals$x2),digits=3)
[1] 5.242
```

- **Range**

```
> round(range(vals$x1),digits=3)
[1] -0.123  9.701
> round(range(vals$x2),digits=3)
[1] -4.291 15.531
```

2. Compute the quantiles for each variable. The quantiles of data set are the 0,25,50,75,and 100 percentiles. Report each to 3 decimal places:

```
> round(quantile(vals$x1),digits=3)
    0%    25%    50%    75%   100%
-0.123  4.083  5.052  5.991  9.701
> round(quantile(vals$x2),digits=3)
    0%    25%    50%    75%   100%
-4.291 -0.132  2.567  9.957 15.531
```

3. Create a histogram (see hist() in R and hist() in Matlab) for each variable using 10 bins. The scale of the y-axis should be in terms of density, not frequncy. Also, overlay a hypothetical true density as a line. For this problem investigate the claim that the data were generated from a normal density with mean (represented by μ) 5 and and standard deviation (represented by σ) of 1.5. The normal probability density function (pdf) is given by:

$$f(x) = \frac{1}{\sqrt{2\pi}\sigma}e^{\frac{-(x-\mu)^2}{2\sigma^2}}$$

This function indicates the y-value for any x-value in the data set. You are asked to plot this line, f(x), and overlay it on each histogram. This plot will be a good way to visually inspect if a variable comes from a normal distribution with the indicated parameters. Note: you may do this anyway you wish, so if there are built in functions that will help you accomplish this, you are free to use them.

```
> hist(vals$x1,freq=F,main="Histogram (Density Based)",xlab="Values of x1",breaks=10)
> y<-exp(-(vals$x1-5)^2/(2*1.5*1.5))/(1.5 * sqrt(2*pi))
> lines(sort(vals$x1),y[order(vals$x1)],col="red")
```
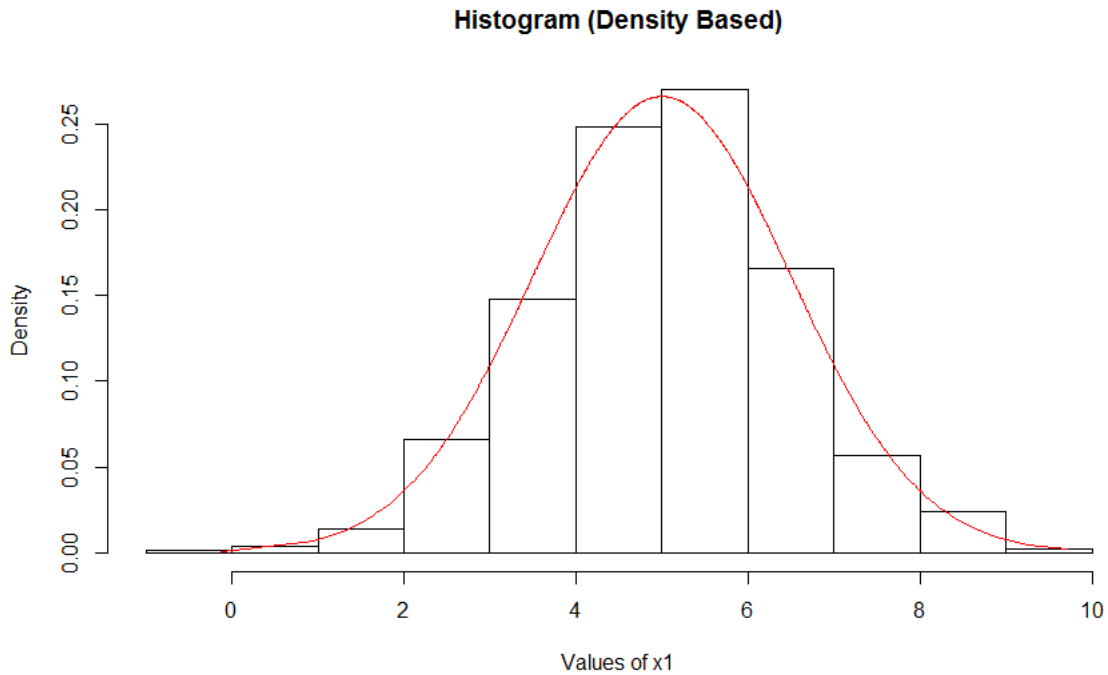
Histogram (Density Based)

Figure 7: The Histogram for values of x1

Replicating the same for the values of x2 we get different results. This is probably because of the reason that the mean for x2 is not 5 and the standard deviation is not 1.5 either. The density() function is plotted in blue to show how the density fucntion would look if superimposed on the graph.

```
> hist(vals$x2,freq=F,main="Histogram (Density Based)",xlab="Values of x2",breaks=10)
> y2<-exp(-(vals$x2-5)^2/(2*1.5*1.5))/(1.5 * sqrt(2*pi))
> lines(sort(vals$x2),y2[order(vals$x2)],col="red")
> lines(density(vals$x2),col="blue")
```
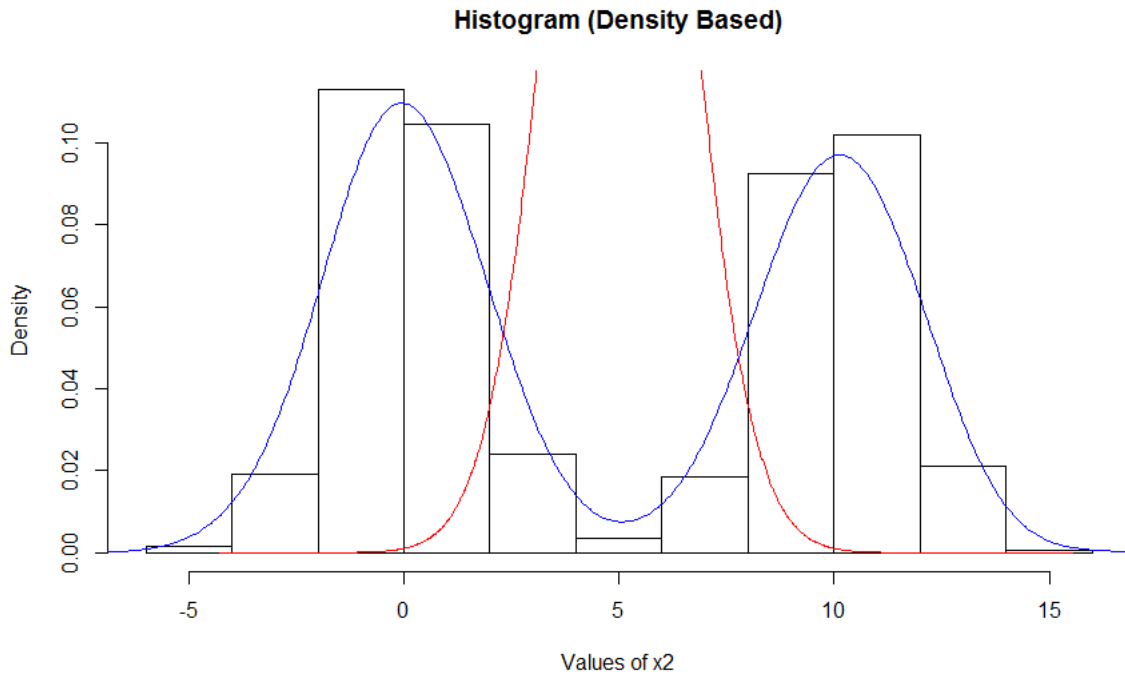
Figure 8: The Histogram for values of x2

4. Create a quantile-quantile plot (commonly called a QQ plot) for each variable. The purpose of a QQ plot is to visually match an observed distribution to a theoretical one. In the previous question, it was suggested that the data came from a normal distribution. If this is true, when we plot the quantiles of our data against the quantiles of a normal distribution, it should form a straight line. Construct your plot using `qqnorm()` in R and `qqplot()` in Matlab. Include in your plot a line indicating perfect agreement, i.e. y = x.

```
> qqnorm(vals$x1,main="QQ Plot for values of x1")
> qqline(vals$x1,col="red",lwd=2)
```
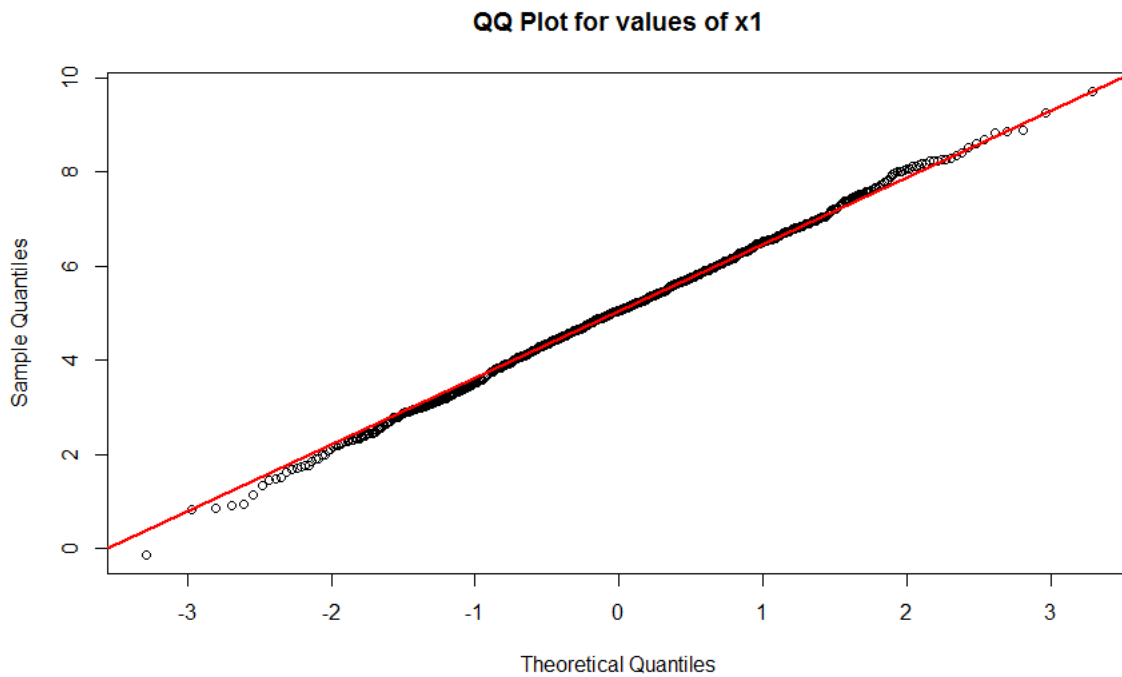
Figure 9: The QQ Plot for values of x1

It is obvious from the following figure that the values of x2 do not form a normal distribution.

```
> qqnorm(vals$x2,main="QQ Plot for values of x2")
> qqline(vals$x2,col="red",lwd=2)
```
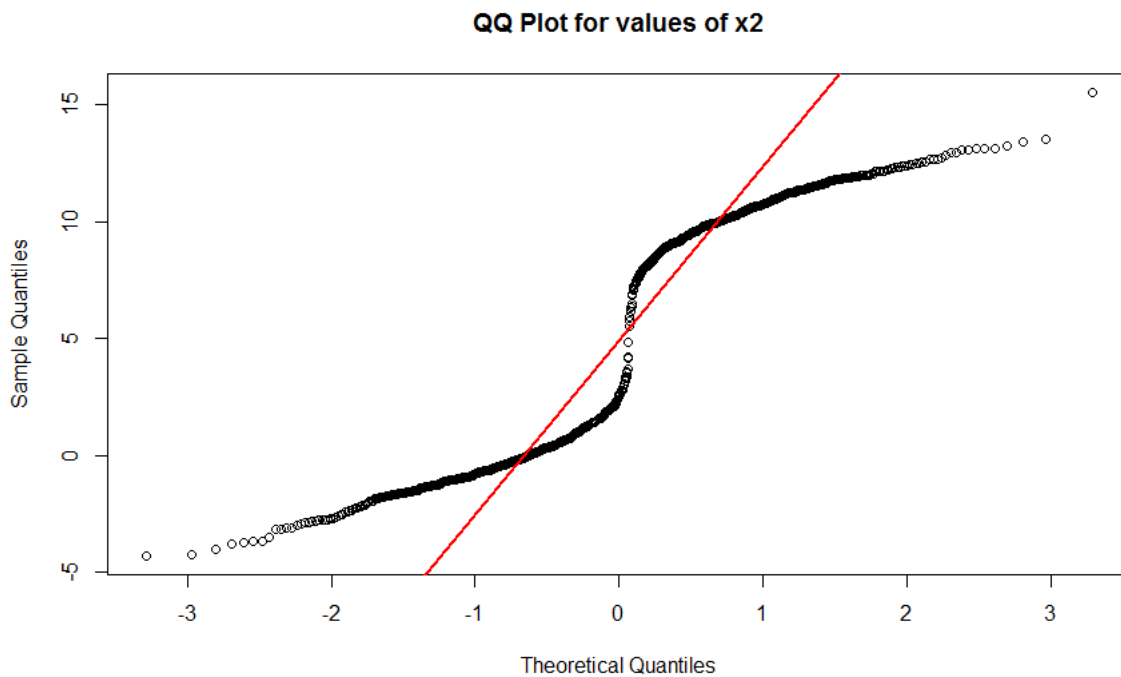

Figure 10: The QQ Plot for values of x2

5. Comment briefly on what you have learned about each variable. Include comparisons between x1 and x2 using location measures (mean, median, etc.), spread measures (standard deviation), and the shape of the histograms. Qualitatively, does either variable appear to have come from a normal distribution? Why?

From the observations done above one can easily conclude that x1 seems to be from a normal distribution while x2 is not.

Let us look at why x1 could be from a normal distribution

- It is symmetric around $x = \mu$ and this can be verified with the help of the histogram.
- The curve and the hsitogram can show that it is unimodal and reaches a maxima at only $x = \mu$
- The QQ plot seems to come up with a roughly linear graph with all the pints roughly falling on the line $y = 4.4x + 5$. This means the distribution is probably normal

. x2 doesnt seem to be normal and the following reasons

- The QQ plot odesnt form a straight line
- The PDF is not unimodal and there are two maxima neither of which are at μ

4 Question 4

For this exercise, you will use the Auto MPG Data Set from the UCI Repository (`http://archive.ics.uci.edu/ml/datasets/Auto+MPG`). Download the auto-mpg.data dataset. DO NOT use auto-mpg.data-original. Once you have downloaded the data, perform the following tasks.

1. Load the dataset

```
> vals<-read.table(file.choose())
> names(vals)<- c("mpg","cylinders","displacement","horsepower",
+ "weight","acceleration","model_yr","origin","car_name")
```

2. Display the first 10 rows of the dataset

```
> head(vals,10)
```

mpg	cylinders	displacement	horsepower	weight	acceleration	model_yr	origin	car_name
18	8	307	130.0	3504	12.0	70	1	chevroletchevellemalibu
15	8	350	165.0	3693	11.5	70	1	buickskylark320
18	8	318	150.0	3436	11.0	70	1	plymouthsatellite
16	8	304	150.0	3433	12.0	70	1	amcrebelsst
17	8	302	140.0	3449	10.5	70	1	fordtorino
15	8	429	198.0	4341	10.0	70	1	fordgalaxie500
14	8	454	220.0	4354	9.0	70	1	chevroletimpala
14	8	440	215.0	4312	8.5	70	1	plymouthfuryiii
14	8	455	225.0	4425	10.0	70	1	pontiaccatalina
15	8	390	190.0	3850	8.5	70	1	amcambassadordpl

3. Let us assume that the dataset is being used for a regression task to predict the MPG (miles per gallon) using x = {Displacement,Horsepower,Weight, Acceleration}

   a) The following linear model was determined empirically to estimate MPG:
   MPGEst = $w_0 + \sum_{i=1}^{n}(w_i x_i)$
   where $w_0 = 45 : 5455$ and $w = \{0.0177, 0.037, 0.0037, 0.2941\}$
   Using the linear model, compute MPGEst for the entire dataset.

```
> vals2<- vals[complete.cases(vals),]
> mpgest <- w0 + ((w[1]*vals2$displacement) +
+ (w[2]*paste(as.numeric(vals2$horsepower))) +
+ (w[3]*vals2$weight) + (w[4]*vals2$acceleration))
```

b) The dataset contains the true value of MPG for each car and the linear model above enabled you to estimate the MPG for each car. Is the MPGEst correct? If not, report the Root Mean Squared Error (RMSE).

$$RMSE = \sqrt{\frac{\sum_{i=0}^{n}(y_i - \hat{y}_i)^2}{n}} \tag{1}$$

where, y corresponds to MPG given in the dataset and $\hat{y}$ corresponds to MPGEst that you computed previously.

```
> > rmse<-sqrt(sum((vals2$mpg-mpgest)^2)/nrow(vals2))
> rmse
[1] 4.407523
```

4. We would like to visualize how MPG varies with Acceleration and Displacement. Show the result of the following:

   a) Display a 3-D scatter plot of Acceleration, Horsepower and MPG. Label the axes.

```
> scatterplot3d(vals2$acceleration,as.numeric(paste(vals2$horsepower)),
+ vals2$mpg,xlab="Acceleration",ylab="Horsepower",zlab="Horse power")
```


Figure 11: The Scatter 3D Plot

   b) In the same 3-D scatter plot, visualize how both MPG and MPGEst vary with reference to Acceleration and Horsepower. Label the axes. Make sure that MPG and MPGEst are displayed in different colors. Show a legend.

```
> pl2<-scatterplot3d(vals2$acceleration,as.numeric(paste(vals2$horsepower),
```

```
+ vals2$mpg,xlab="Acceleration",ylab="Horsepower",
+ zlab="MPG",zlim=c(0,100))
> pl2$points3d(vals2$acceleration,
+ as.numeric(paste(vals2$horsepower)),mpgest,col="red")
> legend("topleft", inset=.05, bty="n", cex=0.8,
+ title="MPG types",  c("MPG Actual", "MPG Estimated"),
+ fill=c("black", "red"))
```


Figure 12: The 3D Plot with estimated mpg