

Topics in Artificial Intelligence:

# Machine Learning Techniques for Data Mining

B Semester 2000

Lecturer: Eibe Frank




# Incremental clustering

- COBWEB/CLASSIT: incrementally forms a hierarchy of clusters
- In the beginning tree consists of empty root node
- Instances are added one by one, and the tree is updated appropriately at each stage
- Updating involves finding the right leaf for an instance (possibly restructuring the tree)
- Updating decisions are based on *category utility*

# Clustering the weather data



ID code	Outlook	Temp.	Humidity	Windy
A	Sunny	Hot	High	False
B	Sunny	Hot	High	True
C	Overcast	Hot	High	False
D	Rainy	Mild	High	False
E	Rainy	Cool	Normal	False
F	Rainy	Cool	Normal	True
G	Overcast	Cool	Normal	True
H	Sunny	Mild	High	False
I	Sunny	Cool	Normal	False
J	Rainy	Mild	Normal	False
K	Sunny	Mild	Normal	True
L	Overcast	Mild	High	True
M	Overcast	Hot	Normal	False
N	Rainy	Mild	High	True

# Steps 1-3



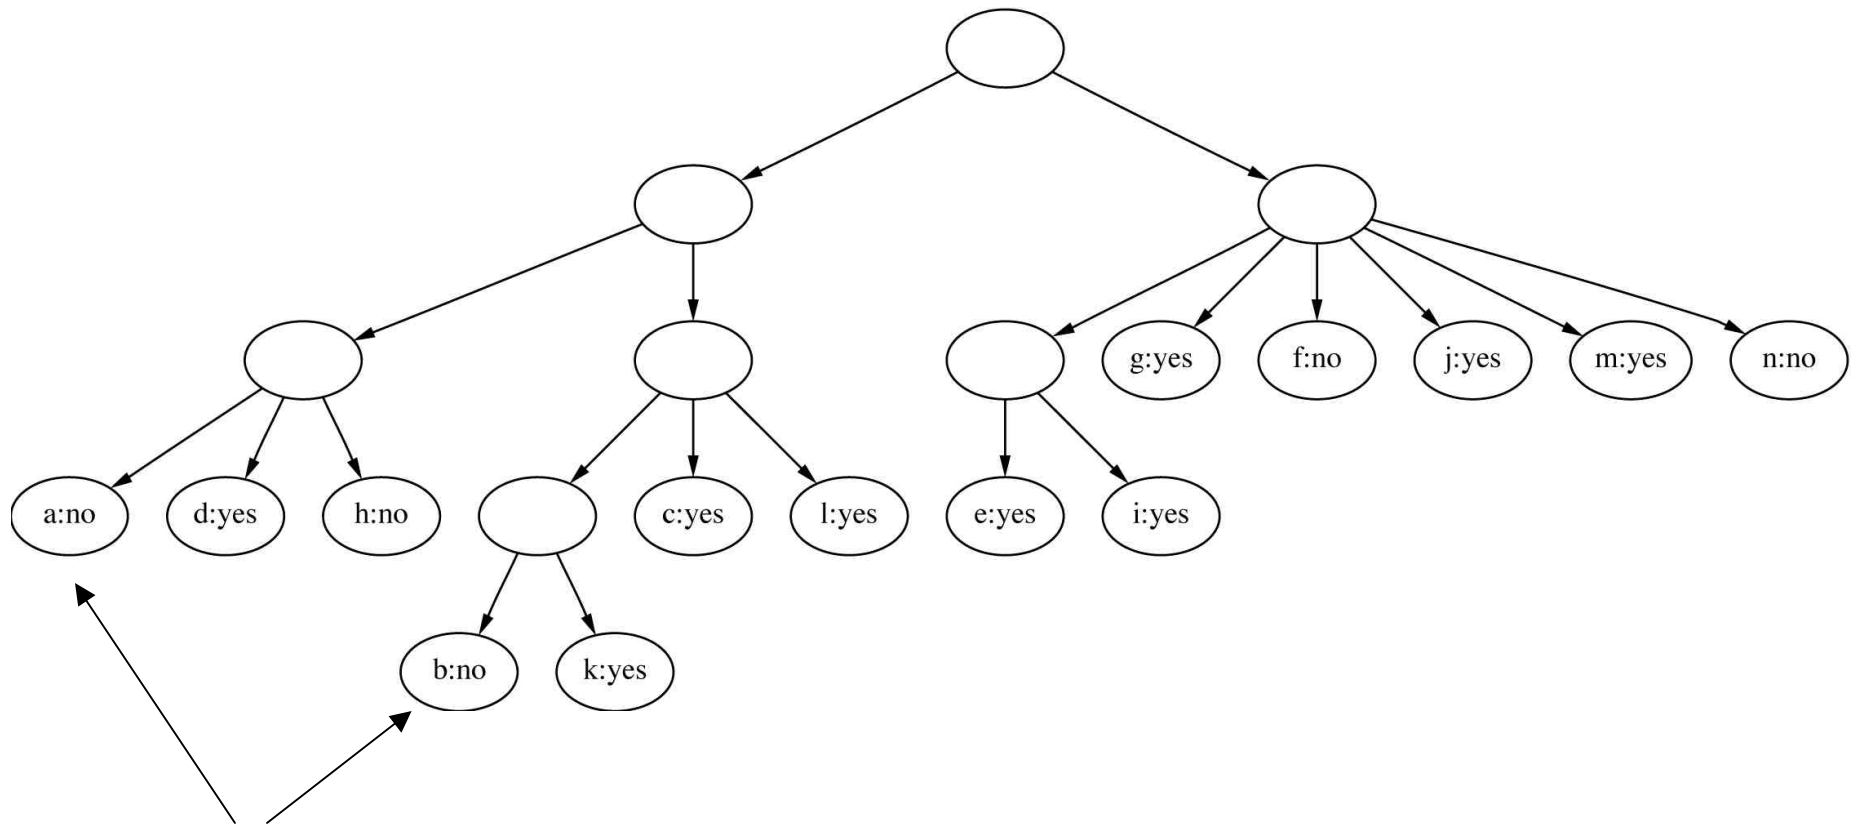
# Steps 3-4

Best host and runner-up have been *merged*



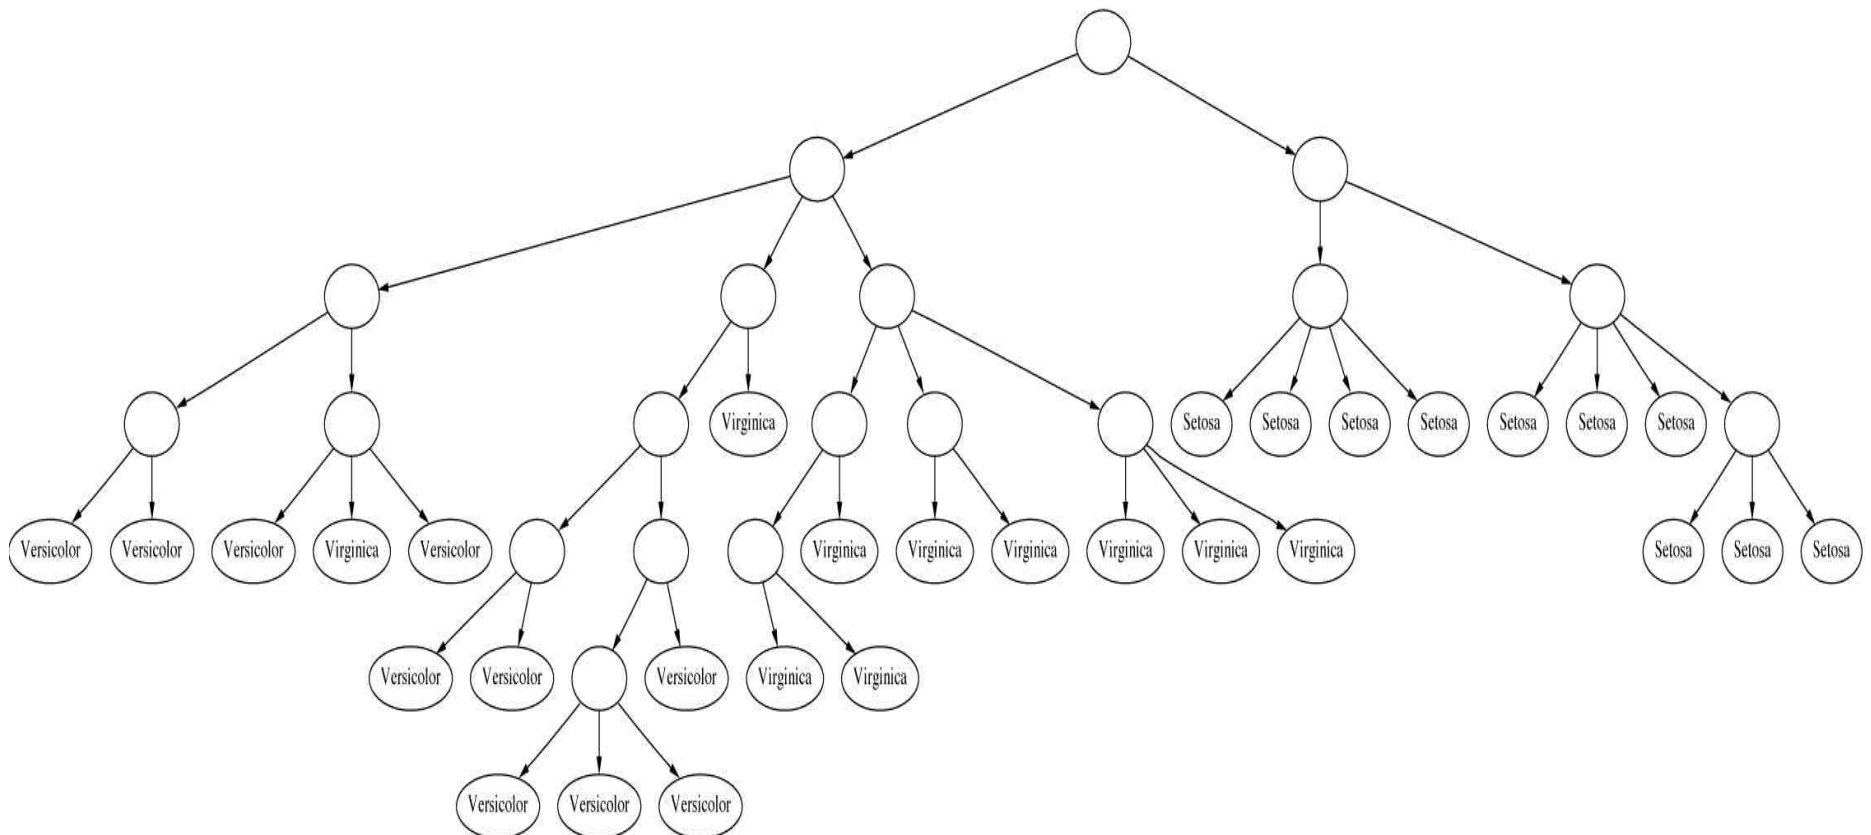
Note: *splitting* the best host is considered if merging doesn't help

# The final hierarchy

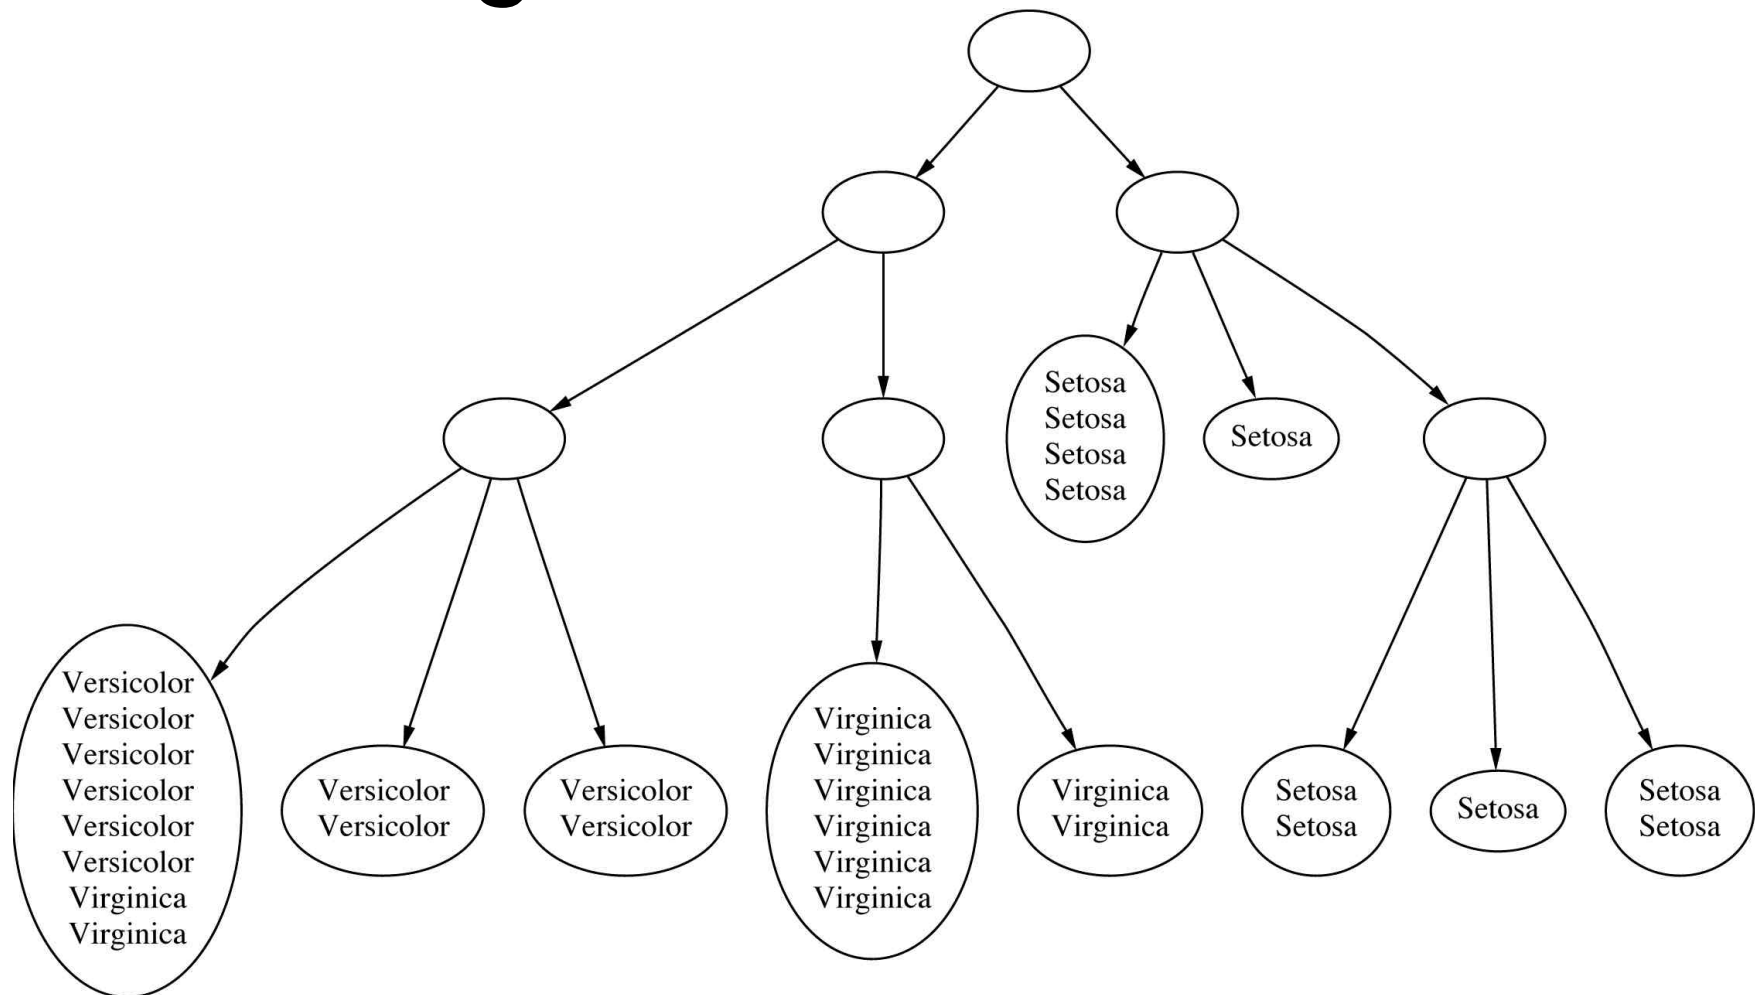


*a and b are actually very similar*

# Clustering (parts) of the iris data



# Clustering the iris data with cutoff





# Category utility

- Category utility is a kind of quadratic loss function defined on conditional probabilities:

$$CU(C_1, C_2, \dots, C_k) = \frac{\sum_l \Pr[C_l] \sum_i \sum_j (\Pr[a_i = v_{ij} | C_l]^2 - \Pr[a_i = v_{ij}]^2)}{k}$$

- If every instance gets put into a different category the numerator becomes ( $m = \text{\#attributes}$ ):

$$m - \Pr[a_i = v_{ij}]^2 \quad \longleftarrow \quad \textit{maximum}$$

# Numeric attributes

- We assume normal distribution:

$$f(a) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(a-\mu)^2}{2\sigma^2}}$$

- Then we get:

$$\sum_j \Pr[a_i = v_{ij}]^2 \Leftrightarrow \int f(a_i)^2 da_i = \frac{1}{2\sqrt{\pi}\sigma_i}$$

- Thus

$$CU = \frac{\sum_l \Pr[C_l] \sum_i \sum_j (\Pr[a_i = v_{ij} | C_l]^2 - \Pr[a_i = v_{ij}]^2)}{k}$$

is

$$CU = \frac{\sum_l \Pr[C_l] \frac{1}{2\sqrt{\pi}} \sum_i \left( \frac{1}{\sigma_{il}} - \frac{1}{\sigma_i} \right)}{k}$$

- *Acuity* parameter: prespecified minimum variance

# Probability-based clustering


- Problems with above heuristic approach:
  - ◆ Division by  $k$ ?
  - ◆ Order of examples?
  - ◆ Are restructuring operations sufficient?
  - ◆ Is result at least *local* minimum of category utility?
- From a probabilistic perspective, we want to find the most likely clusters given the data
- Also: instance only has certain probability of belonging to a particular cluster

# Finite mixtures

- Probabilistic clustering algorithms model the data using a *mixture* of distributions
- Each cluster is represented by one distribution
  - ◆ The distribution governs the probabilities of attributes values in the corresponding cluster
- They are called *finite mixtures* because there is only a finite number of clusters being represented
- Usually individual distributions are normal distribut.
- Distributions are combined using cluster weights

# A two-class mixture model

data					
A	51	B	62	B	64
A	43	A	47	A	51
B	62	A	52	A	51
B	64	B	64	B	63
A	45	A	51	A	43
A	42	B	65	B	65
A	46	A	48	B	66
A	45	A	49	B	65
A	45	A	46	A	46
				A	39
				B	62
				B	64
				A	52
				B	63
				B	64
				A	48
				B	64
				A	48



# Using the mixture model

- The probability of an instance  $x$  belonging to cluster  $A$  is:

$$\Pr[A | x] = \frac{\Pr[x | A] \Pr[A]}{\Pr[x]} = \frac{f(x; \mu_A, \sigma_A) p_A}{\Pr[x]}$$

with

$$f(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

- The *likelihood* of an instance given the clusters is:

$$\Pr[x | \text{the distributions}] = \sum_i \Pr[x | \text{cluster}_i] \Pr[\text{cluster}_i]$$

# Learning the clusters

- Assume we know that there are  $k$  clusters
- To learn the clusters we need to determine their parameters
  - ◆ I.e. their means and standard deviations
- We actually have a performance criterion: the likelihood of the training data given the clusters
- Fortunately, there exists an algorithm that finds a local maximum of the likelihood

# The EM algorithm

- *EM algorithm*: expectation-maximization algorithm
  - ◆ Generalization of *k*-means to probabilistic setting
- Similar iterative procedure:
  1. Calculate cluster probability for each instance (expectation step)
  2. Estimate distribution parameters based on the cluster probabilities (maximization step)
- Cluster probabilities are stored as instance weights



# More on EM

- Estimating parameters from weighted instances:

$$\mu_A = \frac{w_1 x_1 + w_2 x_2 + \dots + w_n x_n}{w_1 + w_2 + \dots + w_n}$$

$$\sigma_A^2 = \frac{w_1 (x_1 - \mu)^2 + w_2 (x_2 - \mu)^2 + \dots + w_n (x_n - \mu)^2}{w_1 + w_2 + \dots + w_n}$$

- Procedure stops when log-likelihood saturates
- Log-likelihood:

$$\sum_i \log(p_A \Pr[x_i | A] + p_B \Pr[x_i | B])$$

# Extending the mixture model

- Using more than two distributions: easy
- Several attributes: easy if independence is assumed
- Correlated attributes: difficult
  - ◆ Modeled jointly using a bivariate normal distribution with a (symmetric) covariance matrix
  - ◆ With  $n$  attributes this requires estimating  $n + n(n+1)/2$  parameters
- Nominal attributes: easy if independent

# More on extensions

- Correlated nominal attributes: difficult
  - ◆ Two correlated attributes result in  $v_1 v_2$  parameters
- Missing values: easy
- Distributions other than the normal distribution can be used:
  - ◆ “log-normal” if predetermined minimum is given
  - ◆ “log-odds” if bounded from above and below
  - ◆ Poisson for attributes that are integer counts
- Cross-validation can be used to estimate  $k$ !!

# Bayesian clustering

- Problem: overfitting possible if number of parameters gets large
- *Bayesian approach*: every parameter has a prior probability distribution
  - ◆ Gets incorporated into the overall likelihood figure and thereby penalizes introduction of parameters
- Example: Laplace estimator for nominal attributes
- Can also have prior on number of clusters!
- Actual implementation: NASA's AUTOCLASS

# Discussion

- Clusters can be interpreted by using supervised learning in a post-processing step
- Can be used to fill in missing values
- May be advantageous to make attributes more independent in pre-processing step
  - ◆ I.e. using *principal component analysis*
- Big advantage of probabilistic clustering schemes:
  - ◆ Likelihood of data can be estimated and used to compare different clustering models